

CLEF Guidelines - Domain Specific 2006

In these Guidelines, we provide information on the CLEF 2006 test collections, the tasks, data manipulation, query construction and results submission for the Domain Specific track.

Test Collection

The Domain Specific 2006 track used test collections of scientific documents:

- GIRT4: a GIRT collection of approx. 150,000 German social science documents. Available partly as two parallel (translation equivalent) corpora in German and English. Controlled vocabularies in German-English and German-Russian are also provided.
- RSSC: Russian Social Science Corpus.
- ISSS database: on economy (99,000 documents) and on sociology (46,000 documents).

25 topics are prepared in English, German and Russian for this track.

Tasks

The 2006 Domain Specific tasks were:

- **Monolingual** task: Query the English, German or Russian social science collections using topics in the same language and submit results in a ranked list.
- **Bilingual** task: The three target collections can be queried using topics in a different language.
- **Multilingual** task: Selecting a topic set in either EN, DE or RU, retrieve relevant documents from the entire social science collection of English, German and Russian documents and submit the results in a single merged and ranked list.

Much of the evaluation methodology is an adaptation of the strategy studied for the TREC Domain Specific task. The instructions given below have been derived from those distributed by TREC.

CONSTRUCTING AND MANIPULATING THE SYSTEM DATA STRUCTURES

The system data structures are defined to consist of the original documents, any new structures built automatically from the documents (such as inverted files, thesauri, conceptual networks, etc.), and any new structures built manually from the documents (such as thesauri, synonym lists, knowledge bases, rules, etc.).

1. The system data structures may not be modified in response to CLEF 2006 topics. For example, you cannot add topic words that are not in your dictionary. The CLEF tasks represent the real-world problem of an ordinary user posing a question to a system. In the case of the cross-language tasks, the question is posed in one language and relevant documents must be retrieved whatever the language in which they have been written. If an ordinary user could not

make the change to the system, you should not make it after receiving the topics.

2. There are several parts of the CLEF data collections that contain manually-assigned, controlled or uncontrolled index terms. These fields are delimited by SGML (XML-compatible) tags. Since the primary focus of CLEF is on retrieval of naturally occurring text over language boundaries, these manually-indexed terms should not be indiscriminately used as if they are a normal part of the text. If your group decides to use these terms, they should be part of a specific experiment that utilizes manual indexing terms, and these runs should be declared as manual runs.

3. Only the following fields may be used for automatic retrieval:

- **GIRT4-DE**: all fields
- **GIRT4-EN**: all fields
- **RSSC**: TITLE, TEXT
- **ISSI**: TITLE-RU, ABSTRACT-RU

Learning from (e.g. building translation sources from) such fields is permissible.

GUIDELINES FOR CONSTRUCTING THE QUERIES

The queries are constructed from the topics. Each topic consists of three fields: a brief title statement; a one-sentence description; a more complex narrative specifying the relevance assessment criteria. Queries can consist of 1 or more of these fields.

There are many possible methods for converting the supplied topics into queries that your system can execute. We have broadly defined two generic methods, "automatic" and "manual", based on whether manual intervention is used or not. When more than one set of results are submitted, the different sets may correspond to different query construction methods, or if desired, can be variants within the same method.

The manual query construction method includes BOTH runs in which the queries are constructed manually and then run without looking at the results AND runs in which the results are used to alter the queries using some manual operation. The distinction is being made here between runs in which there is no human involvement (automatic query construction) and runs in which there is some type of human involvement (manual query construction). It is clear that manual runs should be appropriately motivated in a CLIR context, e.g. a run where a proficient human simply translates the topic into the document language(s) is not what most people think of as cross-language retrieval.

To further clarify this, here are some example query construction methodologies, and their correct query construction classification. Note that these are only examples; many other methods may be used for automatic or manual query construction.

1. queries constructed automatically from the topics, the retrieval results of these queries sent to the CLEF results server --> automatic query construction
2. queries constructed automatically from the topics, then expanded by a method that takes terms automatically from the top 30 documents (no human involved) --> automatic query construction
3. queries constructed manually from the topics, results of these queries sent to the CLEF results server --> manual query construction
4. queries constructed automatically from the topics, then modified by human selection of terms suggested from the top 30 documents --> manual query construction

Note that by including all types of human-involved runs in the manual query construction method we make it harder to do comparisons of work within this query construction method. Therefore groups are strongly encouraged to determine what constitutes a base run for their experiments and to do these runs (officially or unofficially) to allow useful interpretations of the results. For those of you who are new to CLEF, unofficial runs are those not submitted to CLEF but evaluated using the [trec_eval](#) tool created for TREC. (See previous years' CLEF papers for examples of use of base runs.)

WHAT TO DO WITH YOUR RESULTS

Results have to be formatted in ASCII, with one line per document retrieved. The lines have to be formatted as follows:

| Field#1 | Field#2 | Field#3 | Field#4 | Field#5 | Field#6 |
|---------|---------|----------------|---------|----------|----------|
| 10 | Q0 | document.00072 | 0 | 0.017416 | runidex1 |

The fields must be separated by ONE blank and have the following meanings:

1. Query number (eliminate any identifying letters). Please only use SIMPLE numbers ("1", not "001").

INPUT MUST BE SORTED NUMERICALLY BY QUERY NUMBER.

2. Query iteration (will be ignored. Please choose "Q0" for all experiments).
3. Document number (content of the <DOCNO> tag.).
4. Rank 0..n (0 is best matching document. If you retrieve 1000 documents per query, rank will be 0..999, with 0 best and 999 worst). Note that rank starts at 0 (zero) and not 1 (one).

MUST BE SORTED IN INCREASING ORDER PER QUERY.

5. RSV value (system specific value that expresses how relevant your system deems a document to be. This is a floating point value. High relevance should be expressed with a high value). If a document D1 is considered more relevant than a document D2, this must be reflected in the fact that $RSV1 > RSV2$. If $RSV1 = RSV2$, the documents may be randomly reordered during calculation of the evaluation measures. Please use a decimal point ".", not a comma. Do not use any form of separators for thousands. The only legal characters for the RSV values are 0-9 and the decimal point.

MUST BE SORTED IN DECREASING ORDER PER QUERY.

6. Run identifier (please chose an unique ID for each experiment you submit). Only use a-z,

A-Z and 0-9. No special characters, accents, etc.

The fields are separated by a single space.

The file contains nothing but lines formatted in the way described above.

You are expected to retrieve 1000 documents per query. An experiment that retrieves a maximum of 1000 documents each for 20 queries therefore produces a file that contains a maximum of 20000 lines.

You should know that the effectiveness measures used in CLEF evaluate the performance of systems at various points of recall. Participants must thus return at most 1000 documents per query in their results. Please note that by its nature, the average precision measure does not penalize systems that return extra irrelevant documents at the bottom of their result lists.